

## # Image Classification using AWS SageMaker

Thiviyaa S

Use AWS Sagemaker to train a pretrained model that can perform image classification by using the Sagemaker profiling, debugger, hyperparameter tuning and other good ML engineering practices. This can be done on either the provided dog breed classification data set or one of your choice.

### ## Project Set Up and Installation

Enter AWS through the gateway in the course and open SageMaker Studio.

Download the starter files.

Download/Make the dataset available.

### ## Dataset

The provided dataset is the dogbreed classification dataset which can be found in the classroom.

The project is designed to be dataset independent so if there is a dataset that is more interesting or relevant to your work, you are welcome to use it to complete the project.

### ### Access

Upload the data to an S3 bucket through the AWS Gateway so that SageMaker has access to the data.

### ## Hyperparameter Tuning

What kind of model did you choose for this experiment and why? Give an overview of the types of parameters and their ranges used for the hyperparameter search

ResNet50 Model:

-->The ResNet-18 pre-trained model is used in this project.

-->The ResNet-50 is an 18-layer deep convolutional neural network. This network's pre-trained version has been trained on over a million photos from the ImageNet collection. The network can classify photos into 1000 different object categories, including keyboards, mice, pencils, and a variety of animals.

-->The dataset we'll be working with is an Image dataset that contains photographs of dogs.

-->Training, validation, and testing are the three parts of the dataset.

Hyperparameter search:

-->Learning rate : -> default is 0.001 and the chosen range is =[0.0001, 0.1] -> learning rate is a ContinuousParameter.

-->epochs: -> default is 1e-08 and the chosen range is= [1e-9, 1e-8] -> the epochs is a IntegerParameter

-->Weight decay: -> default is 0.01 and the chosen range is = [1e-3, 1e-1]

-->Batch size : -> The chosen range is = [ 64, 128] -> batch-size is a CategoricalParameter

Remember that your README should:

- Include a screenshot of completed training jobs
- Logs metrics during the training process
- Tune at least two hyperparameters
- Retrieve the best best hyperparameters from all your training jobs

Udacity Reviews x AWS Machine Learning Engineer x JupyterLab x ImageClassification-using-AWS x Amazon SageMaker x

https://console.aws.amazon.com/sagemaker/home?region=us-east-1#/hyper-tuning-jobs

Hoststar - Watch TV...

aws Services Search for services, features, blogs, docs, and more [Alt+S] N. Virginia voclabs/user1798275=11172933940 @ 6988-2121-1110

Amazon SageMaker > Hyperparameter tuning jobs

### Hyperparameter tuning jobs

Search hyperparameter tuning jobs

Name	Status	Training completed/total	Creation time	Duration
pytorch-training-220124-1325	Completed	3 / 4	Jan 24, 2022 13:25 UTC	27 minutes

Amazon SageMaker > Hyperparameter tuning jobs

Feedback English (US) © 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Udacity Reviews x AWS Machine Learning Engineer x JupyterLab x ImageClassification-using-AWS x Amazon SageMaker x

https://console.aws.amazon.com/sagemaker/home?region=us-east-1#/models

Hoststar - Watch TV...

aws Services Search for services, features, blogs, docs, and more [Alt+S] N. Virginia voclabs/user1798275=11172933940 @ 6988-2121-1110

Amazon SageMaker > Models

### Models

Search models

Name	ARN	Creation time
pytorch-inference-2022-01-24-14-16-29-054	arn:aws:sagemaker:us-east-1:698821211110:model/pytorch-inference-2022-01-24-14-16-29-054	Jan 24, 2022 14:16 UTC
pytorch-training-2022-01-24-14-06-29-328	arn:aws:sagemaker:us-east-1:698821211110:model/pytorch-training-2022-01-24-14-06-29-328	Jan 24, 2022 14:06 UTC

Amazon SageMaker > Models

Feedback English (US) © 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Udacity Reviews | AWS Machine Learning Engineer | JupyterLab | Amazon SageMaker | ImageClassification-using-AWS

https://console.aws.amazon.com/sagemaker/home?region=us-east-1#/processing-jobs

Hoststar - Watch TV...

aws Services Search for services, features, blogs, docs, and more [Alt+S] N. Virginia voclabs/user1798275=11172933940 @ 6988-2121-1110

Processing jobs

Search processing jobs

	Name	ARN	Creation time	Duration	Status
<input type="radio"/>	pytorch-training-2022-01-2-PoorWeightInitialization-be380358	arn:aws:sagemaker:us-east-1:698821211110:processing-job/pytorch-training-2022-01-2-poorweightinitialization-be380358	Jan 24, 2022 13:53 UTC	14 minutes	Completed
<input type="radio"/>	pytorch-training-2022-01-2-Overtraining-2f5c7f6a	arn:aws:sagemaker:us-east-1:698821211110:processing-job/pytorch-training-2022-01-2-overtraining-2f5c7f6a	Jan 24, 2022 13:53 UTC	10 minutes	Failed
<input type="radio"/>	pytorch-training-2022-01-2-Overfit-f7867fb5	arn:aws:sagemaker:us-east-1:698821211110:processing-job/pytorch-training-2022-01-2-overfit-f7867fb5	Jan 24, 2022 13:53 UTC	12 minutes	Completed
<input type="radio"/>	pytorch-training-2022-01-2-ProfilerReport-12d90181	arn:aws:sagemaker:us-east-1:698821211110:processing-job/pytorch-training-2022-01-2-profilerreport-12d90181	Jan 24, 2022 13:53 UTC	10 minutes	Completed
<input type="radio"/>	pytorch-training-2022-01-2-VanishingGradient-3958ed8b	arn:aws:sagemaker:us-east-1:698821211110:processing-job/pytorch-training-2022-01-2-vanishinggradient-3958ed8b	Jan 24, 2022 13:53 UTC	19 minutes	Completed

Feedback English (US) © 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Udacity Reviews | AWS Machine Learning Engineer | JupyterLab | Amazon SageMaker | ImageClassification-using-AWS

https://console.aws.amazon.com/sagemaker/home?region=us-east-1#/hyper-tuning-jobs/pytorch-training-220124-1325?tab=bestTrainingJob

Hoststar - Watch TV...

aws Services Search for services, features, blogs, docs, and more [Alt+S] N. Virginia voclabs/user1798275=11172933940 @ 6988-2121-1110

Best training job hyperparameters

Search

Name	Type	Value
_tuning_objective_metric	FreeText	average test loss
batch_size	Categorical	"128"
eps	Continuous	1.598855182345577e-09
lr	Continuous	0.00043270131115326176
sagemaker_container_log_level	FreeText	20
sagemaker_estimator_class_name	FreeText	"PyTorch"
sagemaker_estimator_module	FreeText	"sagemaker.pytorch.estimator"
sagemaker_job_name	FreeText	"classification-2022-01-24-13-25-11-297"
sagemaker_program	FreeText	"hpo.py"
sagemaker_region	FreeText	"us-east-1"
sagemaker_submit_directory	FreeText	"s3://sagemaker-us-east-1-698821211110/classification-2022-01-24-13-25-11-297/source/sourcedir.tar.gz"

Feedback English (US) © 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

The top screenshot shows the AWS SageMaker console 'Endpoints' page. The left sidebar lists navigation options: Notebook, Processing, Training, Inference, Edge Manager, and AWS Marketplace. The 'Endpoints' page displays a table of endpoints:

Name	ARN	Creation time	Status	Last updated
pytorch-inference-2022-01-24-14-16-29-382	arn:aws:sagemaker:us-east-1:69882121110:endpoint/pytorch-inference-2022-01-24-14-16-29-382	Jan 24, 2022 14:16 UTC	InService	Jan 24, 2022 14:31 UTC
pytorch-training-2022-01-24-14-06-29-328	arn:aws:sagemaker:us-east-1:69882121110:endpoint/pytorch-training-2022-01-24-14-06-29-328	Jan 24, 2022 14:06 UTC	InService	Jan 24, 2022 14:09 UTC

The bottom screenshot shows the AWS SageMaker console 'Training jobs' page. The left sidebar lists navigation options: Dashboard, Search, SageMaker Domain, Studio, RStudio, Canvas, Images, Ground Truth, Notebook, Processing, Training, Inference, and Edge Manager. The 'Training jobs' page displays a table of training jobs:

Name	Creation time	Duration	Status
pytorch-training-2022-01-24-13-53-30-631	Jan 24, 2022 13:53 UTC	10 minutes	Completed
pytorch-training-220124-1325-004-26fb9d48	Jan 24, 2022 13:37 UTC	12 minutes	Stopped
pytorch-training-220124-1325-003-2d693780	Jan 24, 2022 13:37 UTC	12 minutes	Completed
pytorch-training-220124-1325-002-26e9dd2b	Jan 24, 2022 13:25 UTC	12 minutes	Completed
pytorch-training-220124-1325-001-328eec61	Jan 24, 2022 13:25 UTC	12 minutes	Completed

## ## Debugging and Profiling

**\*\*TODO\*\*:** Give an overview of how you performed model debugging and profiling in Sagemaker

-->The sagemaker.debugger module was used to help with debugging and profiling.

-->The Amazon SageMaker Debugger gives you complete visibility into the training jobs for cutting-edge machine learning models.

-->This SageMaker Debugger module has high-level methods for configuring the Debugger to monitor, profile, and debug your training job.

-->When creating a SageMaker estimator, use the Debugger-specific settings to gain visibility and insights into your training job.

### ### Results

**\*\*TODO\*\***: What are the results/insights did you get by profiling/debugging your model?

**\*\*TODO\*\*** Remember to provide the profiler html/pdf file in your submission.

Provided

### ## Model Deployment

**\*\*TODO\*\***: Give an overview of the deployed model and instructions on how to query the endpoint with a sample input.

**\*\*TODO\*\*** Remember to provide a screenshot of the deployed active endpoint in Sagemaker.

Udacity Reviews

AWS Machine Learning Engineer

JupyterLab

https://d-gu9uafm8wdr5.studio.us-east-1.sagemaker.aws/jupyter/default/lab?

Hotstar - Watch TV...

File Edit View Run Kernel Git Tabs Settings Help

RUNNING INSTANCES

RUNNING APPS

KERNEL SESSIONS

TERMINAL SESSIONS

train\_and\_deploy.ipynb

train\_model.py

README.md

Unknown No Kernel Share

Failed to start kernel

An error occurred (ResourceNotAvailable) when calling the CreateApp operation: App [arn:aws:sagemaker:us-east-1:698821211110:appid-gu9uafm8wdr5/default-1643020794470/kernelgateway/default-science-1-9-ml-t3-medium-1abf340786679689be9d80559395] is currently in [Deleting] status. To restart the App, wait until it is in [Deleted] or [Failed] status.

Model Deploying

[49]: # TODO: Deploy your model to an endpoint  
predictor=estimator.deploy(initial\_instance\_count=1, instance\_type="ml.m5.xlarge")  
  
-----  
[50]: model\_data = estimator.output\_path + estimator.latest\_training\_job.job\_name + "/output/model.tar.gz"  
print(f"Model: {model\_data}")  
Model: s3://sagemaker-us-east-1-698821211110/pytorch-training-2022-01-24-13-53-30-631/output/model.tar.gz  
[52]: from sagemaker.pytorch import PyTorchModel  
from sagemaker.predictor import Predictor  
jpeg\_serializer = sagemaker.serializers.IdentitySerializer("image/jpeg")  
json\_deserializer = sagemaker.deserializers.JSONDeserializer()  
  
class ImgPredictor(Predictor):  
 def \_\_init\_\_(self, endpoint\_name, sagemaker\_session):  
 super(ImgPredictor, self).\_\_init\_\_(  
 endpoint\_name,  
 sagemaker\_session = sagemaker\_session,  
 serializer = jpeg\_serializer,  
 deserializer = json\_deserializer  
 )  
  
pytorch\_model = PyTorchModel(model\_data = model\_data,  
 role = role,  
 entry\_point = "inference.py",  
 py\_version = "py36",  
 framework\_version = "1.6",  
 predictor\_cls = ImgPredictor  
 )  
  
predictor = pytorch\_model.deploy(initial\_instance\_count = 1, instance\_type = "ml.t2.medium")  
-----  
[54]: from PIL import Image  
import io  
import os  
import numpy as np  
  
test\_dir = "./testImages"  
test\_images = [ 'dog1.jpg', 'dog2.jpg', 'dog3.jpg' ]  
test\_images\_expected\_output = [129, 5, 21]  
for index in range(len(test\_images)):

0 0 Git: refreshing... No Kernel | Starting Kernel: Starting... | Instance MEM Mode: Command Ln 1, Col 1 train\_and\_deploy.ipynb 15:02 25-01-2022

Udacity Reviews

AWS Machine Learning Engineer

JupyterLab

https://d-gu9uafm8wdr5.studio.us-east-1.sagemaker.aws/jupyter/default/lab?

Hotstar - Watch TV...

File Edit View Run Kernel Git Tabs Settings Help

RUNNING INSTANCES

RUNNING APPS

KERNEL SESSIONS

TERMINAL SESSIONS

train\_and\_deploy.ipynb

train\_model.py

README.md

Unknown No Kernel Share

Failed to start kernel

An error occurred (ResourceNotAvailable) when calling the CreateApp operation: App [arn:aws:sagemaker:us-east-1:698821211110:appid-gu9uafm8wdr5/default-1643020794470/kernelgateway/default-science-1-9-ml-t3-medium-1abf340786679689be9d80559395] is currently in [Deleting] status. To restart the App, wait until it is in [Deleted] or [Failed] status.

super(ImgPredictor, self).\_\_init\_\_(  
 endpoint\_name,  
 sagemaker\_session = sagemaker\_session,  
 serializer = jpeg\_serializer,  
 deserializer = json\_deserializer  
)

pytorch\_model = PyTorchModel(model\_data = model\_data,  
 role = role,  
 entry\_point = "inference.py",  
 py\_version = "py36",  
 framework\_version = "1.6",  
 predictor\_cls = ImgPredictor  
 )  
  
predictor = pytorch\_model.deploy(initial\_instance\_count = 1, instance\_type = "ml.t2.medium")  
-----

[54]: from PIL import Image  
import io  
import os  
import numpy as np  
  
test\_dir = "./testImages"  
test\_images = [ 'dog1.jpg', 'dog2.jpg', 'dog3.jpg' ]  
test\_images\_expected\_output = [129, 5, 21]  
for index in range(len(test\_images)):

0 0 Git: refreshing... No Kernel | Starting Kernel: Starting... | Instance MEM Mode: Command Ln 1, Col 1 train\_and\_deploy.ipynb 15:02 25-01-2022

Udacity Reviews | AWS Machine Learning Engineer | JupyterLab

https://d-gu9uafm8wdrs.studio.us-east-1.sagemaker.aws/jupyter/default/lab?

Hotstar - Watch TV...

File Edit View Run Kernel Git Tabs Settings Help

Feedback

RUNNING INSTANCES

RUNNING APPS

KERNEL SESSIONS

TERMINAL SESSIONS

train\_and\_deploy.ipynb | train\_model.py | README.md

Unknown No Kernel Share

Failed to start kernel

An error occurred (ResourceInUseException) when calling the CreateApp operation: App [arn:aws:sagemaker:us-east-1:698521211110:appid-gu9uafm8wdrs/default-1643020704470/kernelgateway/datascience-1.0.ml.t3.medium.1abf34076676960be9d6559395] is currently in [Deleting] status. To restart the App, wait until it is in Deleted or Failed status.

Completed IRALIN plot  
Completed EVAL plot

Steps (TRAIN)	CrossEntropyLoss_output_0 (TRAIN)	Steps (EVAL)	CrossEntropyLoss_output_0 (EVAL)
0	5.0	0	2.7
20	4.2	2	2.5
40	3.4	4	2.3
60	2.6	6	2.2
80	1.8	8	2.1
100	2.1	10	2.1

Git: refreshing... No Kernel | Starting Kernel: Starting... | Instance MEM

Mode: Command | Ln 1, Col 1 | train\_and\_deploy.ipynb

ENG IN 15:03 25-01-2022